

VISUOTACTILE WORLD-MODEL DISAGREEMENT: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study multimodal robot world models. The motivating bet is: Use persistent vision-touch disagreement as a trigger for belief repair and exploration.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For multimodal robot world models, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a; 2025a; 2026b; 2025b;c; 2024). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.530	0.470
observed_only_wam	hidden_mode_shift	0.175	0.825
augmentation_as_data_wam	nominal	0.695	0.305
augmentation_as_data_wam	hidden_mode_shift	0.260	0.740
scalar_uncertainty_wam	nominal	0.667	0.333
scalar_uncertainty_wam	hidden_mode_shift	0.343	0.657
proposed_branch_mechanism	nominal	0.767	0.233
proposed_branch_mechanism	hidden_mode_shift	0.568	0.432

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Visuotactile World-Model Disagreement argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Omdet: Large-scale vision-language multi-dataset pre-training with multimodal detection network, 2024. <https://doi.org/10.1049/cvi2.12268>.
- Mirsam: multimodal vision-language segment anything model for infrared small target detection, 2025a. <https://doi.org/10.1007/s44267-025-00075-0>.
- Everydayvla: A vision-language-action model for affordable robotic manipulation, 2025b. <http://arxiv.org/abs/2511.05397v1>.
- Vla-mp: A vision-language-action framework for multimodal perception and physics-constrained action generation in autonomous driving, 2025c. <https://doi.org/10.3390/s25196163>.
- Multimodal ai: Palm-e’s role in vision-language-robotics & the future of efficient fine-tuning, 2026a. <https://doi.org/10.66967/jaics.2026.v2i103>.
- Vision language models in healthcare through a multimodal approach to medical imaging and clinical applications, 2026b. <https://doi.org/10.4018/979-8-3373-7315-7.ch009>.